

the trusted technology learning source

[Home](#) > [Articles](#) > [Data](#)

Introduction to Polyglot Persistence: Using Different Data Storage Technologies for Varying Data Storage Needs




By [Pramod J. Sadalage](#) and [Martin Fowler](#)
Sep 5, 2012

[Contents](#) [Print](#) [Share This](#) [Discuss](#)

[< Back](#) [Page 5 of 8](#) [Next >](#)

This chapter is from the book



[NoSQL Distilled: A Brief Guide to the Emerging World of Polyglot Persistence](#)

[Learn More](#) [Buy](#)

13.5. Choosing the Right Technology

There is a rich choice of data storage solutions. Initially, the pendulum had shifted from speciality databases to a single RDBMS database which allows all types of data models to be stored, although with some abstraction. The trend is now shifting back to using the data storage that supports the implementation of solutions natively.


If we want to recommend products to customers based on what's in their shopping carts and which other products were bought by customers who bought those products, it can be implemented in any of the data stores by persisting the data with the correct attributes to answer our questions. The trick is to use the right technology, so that when the questions change, they can still be asked with the same data store without losing existing data or changing it into new formats.

Related Resources

[Store](#)

[Articles](#)

[Blogs](#)



[NoSQL for Mere Mortals](#)

By [Dan Sullivan](#)

Book \$39.99



[Using Person-Centered Health Analytics to Live Longer: Leveraging Engagement, Behavior Change, and Technology for a Healthy Life](#)

By [Dwight McNeill](#)

Book \$63.99



[Trends and Research in the Decision Sciences: Best Papers from the 2014 Annual Conference](#)

By [Decision Sciences Institute](#), [Merrill Warkentin](#)

Book \$67.99

[See All Related Store Items](#)

Let's go back to our new feature need. We can use RDBMS to solve this using a hierarchical query and modeling the tables accordingly. When we need to change the traversal, we will have to refactor the database, migrate the data, and start persisting new data. Instead, if we had used a data store that tracks relations between nodes, we could have just programmed the new relations and keep using the same data store with minimal changes.

[+ Share This](#) [📌 Save To Your Account](#)

[< Back](#) [Page 5 of 8](#) [Next >](#)

Discussions

Comments for this thread are now closed.



Comments

Community

1

Login ▾

♥ Recommend

Sort by Oldest ▾

This discussion has been closed.